

State predicting Nim Agent

Florentin-Cristian Udrea (S319029)

Impartial game theory

- Players taking turns
- No hidden information
- Players have the same legal moves from any given position

Impartial game theory

- Players taking turns
- No hidden information
- Players have the same legal moves from any given position




The position you are in defines if you're winning or losing

P-positions & N-positions

Player 1


moves



Player 2

thinks
(hopefully)

moves




[...]

P-positions & N-positions

Player 1


moves



Player 2

thinks
(hopefully)


moves



[...]

Every state can be defined as
a P-position or a N-position

P-positions & N-positions



N-position: the game is in a n-position if a win is secured for the **N**ext player to move

P-position: the game is in a n-position if a win is secured for the **P**revious player that moved

Idea & Strategy



Always choose the move that gets the next state in a P-position



Idea & Strategy



Always choose the move that gets the next state in a P-position



Use a binary classifier to predict if the new state is a P-position or N-position

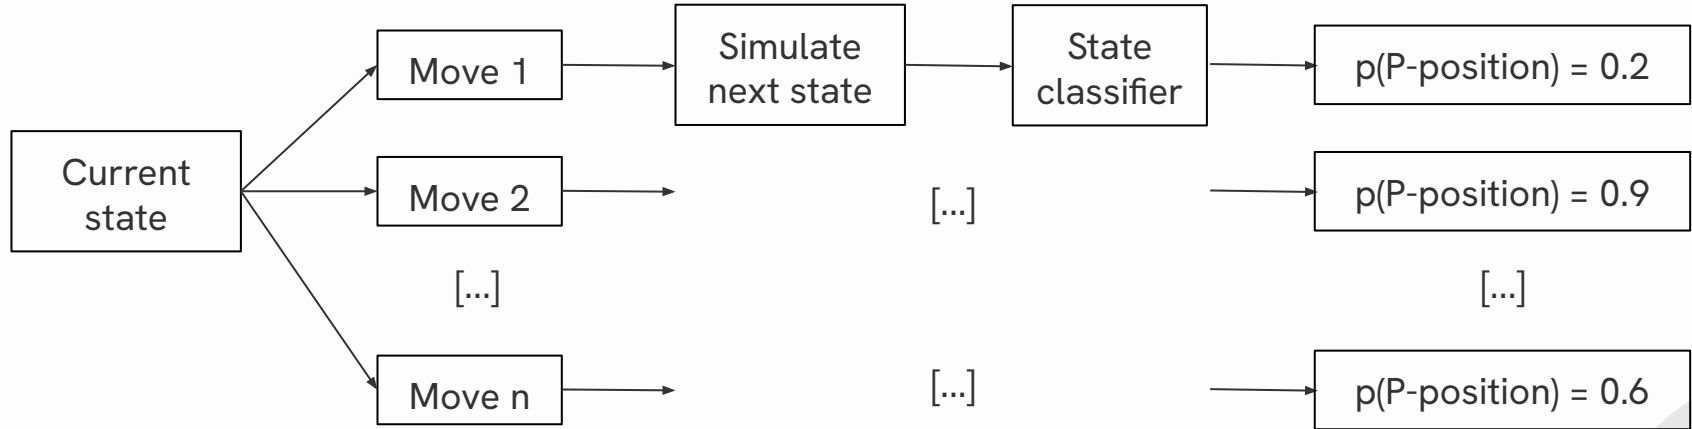
Idea & Strategy



Always choose the move that gets the next state in a P-position




Use a binary classifier to predict if the new state is a P-position or N-position



Regression for Classification - Sigmoid


Linear regression, then put into the sigmoid function to get a (0,1) bounded result



Regression for Classification - Sigmoid

Linear regression, then put into the sigmoid function to get a (0,1) bounded result

Can be interpreted as binomial probability




Regression for Classification - Sigmoid

Linear regression, then put into the sigmoid function to get a (0,1) bounded result

Can be interpreted as binomial probability

Mainly known as '*Logistic classifier*'


```
def get_logistic_score(features, weights):  
    z = np.dot(features, weights)  
    logistic_score = 1 / ( 1 + np.exp(-z) )  
    return logistic_score
```



What are we predicting from ?

General, hand-extracted game features, such as:


- Number of non-zero rows remaining



What are we predicting from ?

General, hand-extracted game features, such as:


- Number of non-zero rows remaining
- Number of non-zero rows remaining having odd number of objects



What are we predicting from ?

General, hand-extracted game features, such as:


- Number of non-zero rows remaining
- Number of non-zero rows remaining having odd number of objects
- Maximum objects on a row



What are we predicting from ?

General, hand-extracted game features, such as:

- Number of non-zero rows remaining
- Number of non-zero rows remaining having odd number of objects
- Maximum objects on a row
- Mean value of non zero rows
- Etc.




What are we predicting from ?

In order to grasp nonlinear relationships I added their polynomial combinations up to the **third degree**, for a total of 165 features


$$[x_1, x_2] \longrightarrow [1, x_1, x_2, x_1^2, x_1x_2, x_2^2]$$

(example: only up to the second degree, having only two features)



No supervised learning

- Classification tasks use annotated data
- We have no data or annotations




No supervised learning

- Classification tasks use annotated data
- We have no data or annotations



Classification performance:
Indirectly estimate using the
average win-rate as a measure



No supervised learning


- Classification tasks use annotated data
- We have no data or annotations



Classification performance:
Indirectly estimate using the average win-rate as a measure



Optimization strategy:
Use evolutionary strategy to improve the weights of the regressor



Improvise. Adapt. Overcome


Implementation details

- **Mutation:** one random weight gets changed by a gaussian sampled value
- **Standard deviation:** gets lower linearly as we get to the final generations (assuming we explored enough and want to fine-tune)
- **Cross-over:** one-cut crossover
- **Selective pressure:** parents are selected randomly with probabilities proportional to their fitness (only 25% of the population)

Even more details

- *Population size* = 30
- *Mutation rate* = 0.2
- *Number of generations* = 150
- *Fitness matches* = 100

Results



(*) made after lab deadline, so not on github yet

Final agent trained against *pure_random* in fitness

```
Fitness (vs. random) -> 0.893  
Fitness (vs. gabriele) -> 1.0  
Fitness (vs. optimal) -> 0.656
```

→ average winrate
over 1000 games

Final agent trained against *optimal* in fitness (*)

```
Fitness (vs. random) -> 0.926  
Fitness (vs. gabriele) -> 1.0  
Fitness (vs. optimal) -> 0.969
```

→ average winrate
over 1000 games

Issues

- Very slow training (3.5h for 150 generations, with multithreading)

Issues


- Very slow training (3.5h for 150 generations, with multithreading)
- Has to evaluate ALL possible moves, may be slow if Nim dimension is high

Issues

- Very slow training (3.5h for 150 generations, with multithreading)
- Has to evaluate ALL possible moves, may be slow if Nim dimension is high
- Result depends on the agent used in the fitness

Issues

- Very slow training (3.5h for 150 generations, with multithreading)
- Has to evaluate ALL possible moves, may be slow if Nim dimension is high
- Result depends on the agent used in the fitness
- Very dependent on the quality of the features



Thanks!

Florentin-Cristian Udrea (S319029)

